

ASIO & CORO

Network Programming with C++20 Coroutines

Part I: From **sync** to **async** to **as-sync**

Peter Eisenlohr
peter@eisenlohr.org
github.com/pgit/asio-coro



OVERVIEW

- Introduction
- Part I: From **sync** to **async** to **as-sync**
- Part II: Working with Coroutines
- Epilogue

INTRODUCTION

ASIO

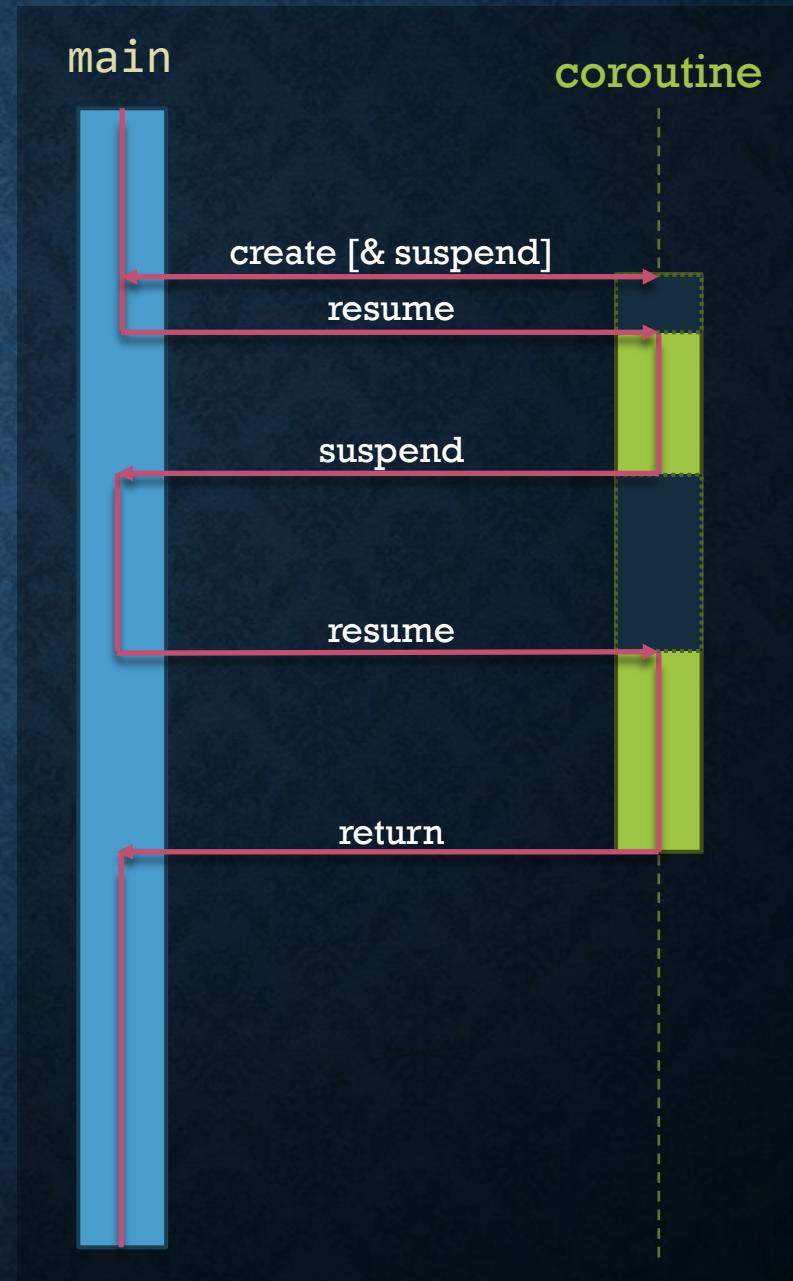
- ASIO (Asynchronous Input/Output) is a cross-platform C++ library for network and low-level I/O programming. It provides a consistent asynchronous model using modern C++.
- P2444 The Asio asynchronous model, Christopher Kohlhoff, 2021
<https://www.open-std.org/jtc1/sc22/wg21/docs/papers/2021/p2444r0.pdf>
- “Powerful but Complex”
- Lowlevel (e.g., no HTTP) – but there are other libraries that build on top of it
Boost::{Process, Beast, MySQL, Redis, MQTT5}, asio-grpc, ...

ASIO

- Key Features:
 - Part of Boost, but also available standalone
 - Supports synchronous and asynchronous operations
 - Works with sockets, timers, serial ports, pipes and more
 - Integrates naturally with C++ coroutines
- Why Use It?
 - Efficient event-driven I/O
 - Scales well for high-performance servers and clients
 - Clean abstraction over platform-specific APIs (epoll, IOCP, etc.)

COROUTINES

- A C++20 coroutine is a special kind of function that can **suspend** and **resume** its execution without blocking the thread.
- New keywords: `co_await`, `co_yield`, and `co_return`.
- Transformed into a state machine by the compiler.
- Enables asynchronous and lazy evaluation patterns.
- Ideal for networking (e.g. with Boost.Asio) where non-blocking IO is crucial.
- **DISCLAIMER:**
This presentation is NOT about the C++20 Coroutines Language Feature.



?

PART I
FROM SYNC TO ASYNC TO AS-SYNC

FROM SYNC TO ASYNC TO AS-SYNC

- Sync Echo Server
- Async Echo Server
- Completion Handlers and Tokens
- The `deferred` completion token
- Coro Echo Server
- `co_spawn`

HELLO WORLD

A minimal **sync** Echo Server

github.com/pgit/asio-coroutine

```
void session(tcp::socket socket)
{
    std::array<char, 64*1024> data;
    for (;;)
    {
        boost::system::error_code ec;
        size_t n = socket.read_some(buffer(data), ec);
        if (ec == error::eof)
            return;
        write(socket, buffer(data, n));
    }
}

void server(tcp::acceptor acceptor)
{
    for (;;)
        session(acceptor.accept());
}

int main()
{
    io_context context;
    server(tcp::acceptor{context, {tcp::v6(), 55555}});
}
```

```
#include <boost/asio.hpp>
using namespace boost::asio;
using ip::tcp;
```

What's wrong
here?

```
void session(tcp::socket socket)                                #include <boost/asio.hpp>
{
    std::array<char, 64*1024> data;
    for (;;)
    {
        boost::system::error_code ec;
        size_t n = socket.read_some(buffer(data), ec);
        if (ec == error::eof)
            return;
        write(socket, buffer(data, n));
    }
}

void server(tcp::acceptor acceptor)
{
    for (;;)
        std::thread(session, acceptor.accept()).detach(); // was: session(acceptor.accept());
}

int main()
{
    io_context context;
    server(tcp::acceptor{context, {tcp::v6(), 55555}});
}
```

SYNCHRONOUS OPERATIONS

PROS & CONS

- Natural control flow (`for`, `if`, `while`, ...)
- Can use (blocking) child functions
- Temporary storage is managed naturally (RAII)
- Lifetime of arguments is clear, including the ability to safely pass temporaries
- Allows for Structured Programming within the thread
- Slow – a context switch between threads takes thousands of cycles
- Does not scale well with many connections
- Requires extra synchronization when accessing shared resources
- Cancellation is difficult

```
void session(tcp::socket socket)
{
    std::array<char, 64*1024> data;
    for (;;)
    {
        boost::system::error_code ec;
        size_t n = socket.read_some(buffer(data), ec);
        if (ec == error::eof)
            return;
        write(socket, buffer(data, n));
    }
}

void server(tcp::acceptor acceptor)
{
    for (;;)
        std::thread(session, acceptor.accept()).detach();
}

int main()
{
    io_context context;
    server(tcp::acceptor{context, {tcp::v6(), 55555}});
}
```

```
#include <boost/asio.hpp>
using namespace boost::asio;
using ip::tcp;
```

?

HELLO ASYNC WORLD

A minimal **async** Echo Server

github.com/pgit/asio-coro

```
class server;                                #include <boost/asio.hpp>

using boost::system::error_code;
using namespace boost::asio;
using ip::tcp;
```



```
int main()
{
    io_context context;
    server server(tcp::acceptor{context, {tcp::v6(), 55555}});
}
```

```
class server
{
public:
    explicit server(tcp::acceptor acceptor)
        : acceptor_(std::move(acceptor))
        do_accept();
    }

private:
    void do_accept()
    {
        acceptor_.async_accept([this](error_code ec, tcp::socket socket)
        {
            do_accept();
        });
    }

    tcp::acceptor acceptor_;
};
```

Control flow:

- sync —————→
- async→

WHAT DID WE JUST DO?

- Unlike `accept` before, `async_accept` returns immediately
- Completion Handler is invoked when operation has completed
- We are no longer in control of the program's execution
- Instead, the framework calls us when something happens
- „Hollywood Principle“ – We call you, don't call us!
- This is „Inversion of Control“:
- `main()` needs to go into `context.run()`

```
#include <boost/asio.hpp>

using boost::system::error_code;
using namespace boost::asio;
using ip::tcp;

int main()
{
    io_context context;
    server server(tcp::acceptor{context, {tcp::v6(), 55555}});
    context.run();
}
```

?

```
class server
{
public:
    explicit server(tcp::acceptor acceptor)
        : acceptor_(std::move(acceptor))
        do_accept();
}

private:
    void do_accept()
    {
        acceptor_.async_accept([this](error_code ec, tcp::socket socket)
            do_accept();
        });
    }

    tcp::acceptor acceptor_;
};
```

```
class server
{
public:
    explicit server(tcp::acceptor acceptor)
        : acceptor_(std::move(acceptor))
        do_accept();
}

private:
    void do_accept()
    {
        acceptor_.async_accept([this](error_code ec, tcp::socket socket)
            if (!ec)
                std::make_shared<session>(std::move(socket))->start();
            do_accept();
        });
    }

    tcp::acceptor acceptor_;
};
```

```
class session : public std::enable_shared_from_this<session>
{
public:
    session(tcp::socket socket) : socket_(std::move(socket)) {}
    void start() { do_read(); }

private:
    void do_read()
    {
        auto self(shared_from_this());
        socket_.async_read_some(buffer(data_),
                               [this, self](error_code ec, size_t length)
        {
            if (!ec)
                do_write(length);
        });
    }

    void do_write(size_t length) { ... }

    tcp::socket socket_;
    std::array<uint8_t, 64*1024> data_;
};
```

```
void do_read()
{
    auto self(shared_from_this());
    socket_.async_read_some(buffer(data_),
                           [this, self](error_code ec, size_t length)
                           {
                               if (!ec)
                                   do_write(length);
                           });
}

void do_write(size_t length)
{
    auto self(shared_from_this());
    async_write(socket_, buffer(data_, length),
               [this, self](error_code ec, size_t /* length */)
               {
                   if (!ec)
                       do_read();
               });
}
```

Control flow:

- sync →
- async ⏪

State:

- socket_
- data_

HELLO ASYNC WORLD

- Global event loop object: `io_context`
- `run()` function, keeps running the IO context as long as there is work
- Server:
 - initiates `async_accept()`. On completion, spawns a new `session`
- Session:
 - initiates `async_read_some()`, completes with some bytes read
 - initiates `async_write()`, completes after all bytes have been written
 - uses state (`socket_`, `data_`) stored in session class
 - uses `shared_ptr<>` to keep session alive



RECAP: SYNCHRONOUS OPERATIONS

PROS & CONS

- Natural control flow (`for`, `if`, `while`, ...)
- Can use (blocking) child functions
- Temporary storage is managed naturally (RAII)
- Lifetime of arguments is clear, including the ability to safely pass temporaries
- Allow for Structured Programming
- Slow – a context switch between threads takes thousands of cycles
- Does not scale well with many connections
- Requires extra synchronization when accessing shared resources
- Cancellation is difficult

ASYNCHRONOUS OPERATIONS

PROS & CONS

- Fast – a scheduler designed specifically for events is 1-2 magnitudes faster
- Scales to thousands of connections
- No synchronization required – single thread only
- ✗ Cancellation is still difficult
- Natural control flow (~~for, if, while, ...~~)
- Can use (blocking) child functions
- Temporary storage is managed naturally (RAII)
- Lifetime of arguments is clear, including the ability to safely pass temporaries
- Allow for Structured Programming

?

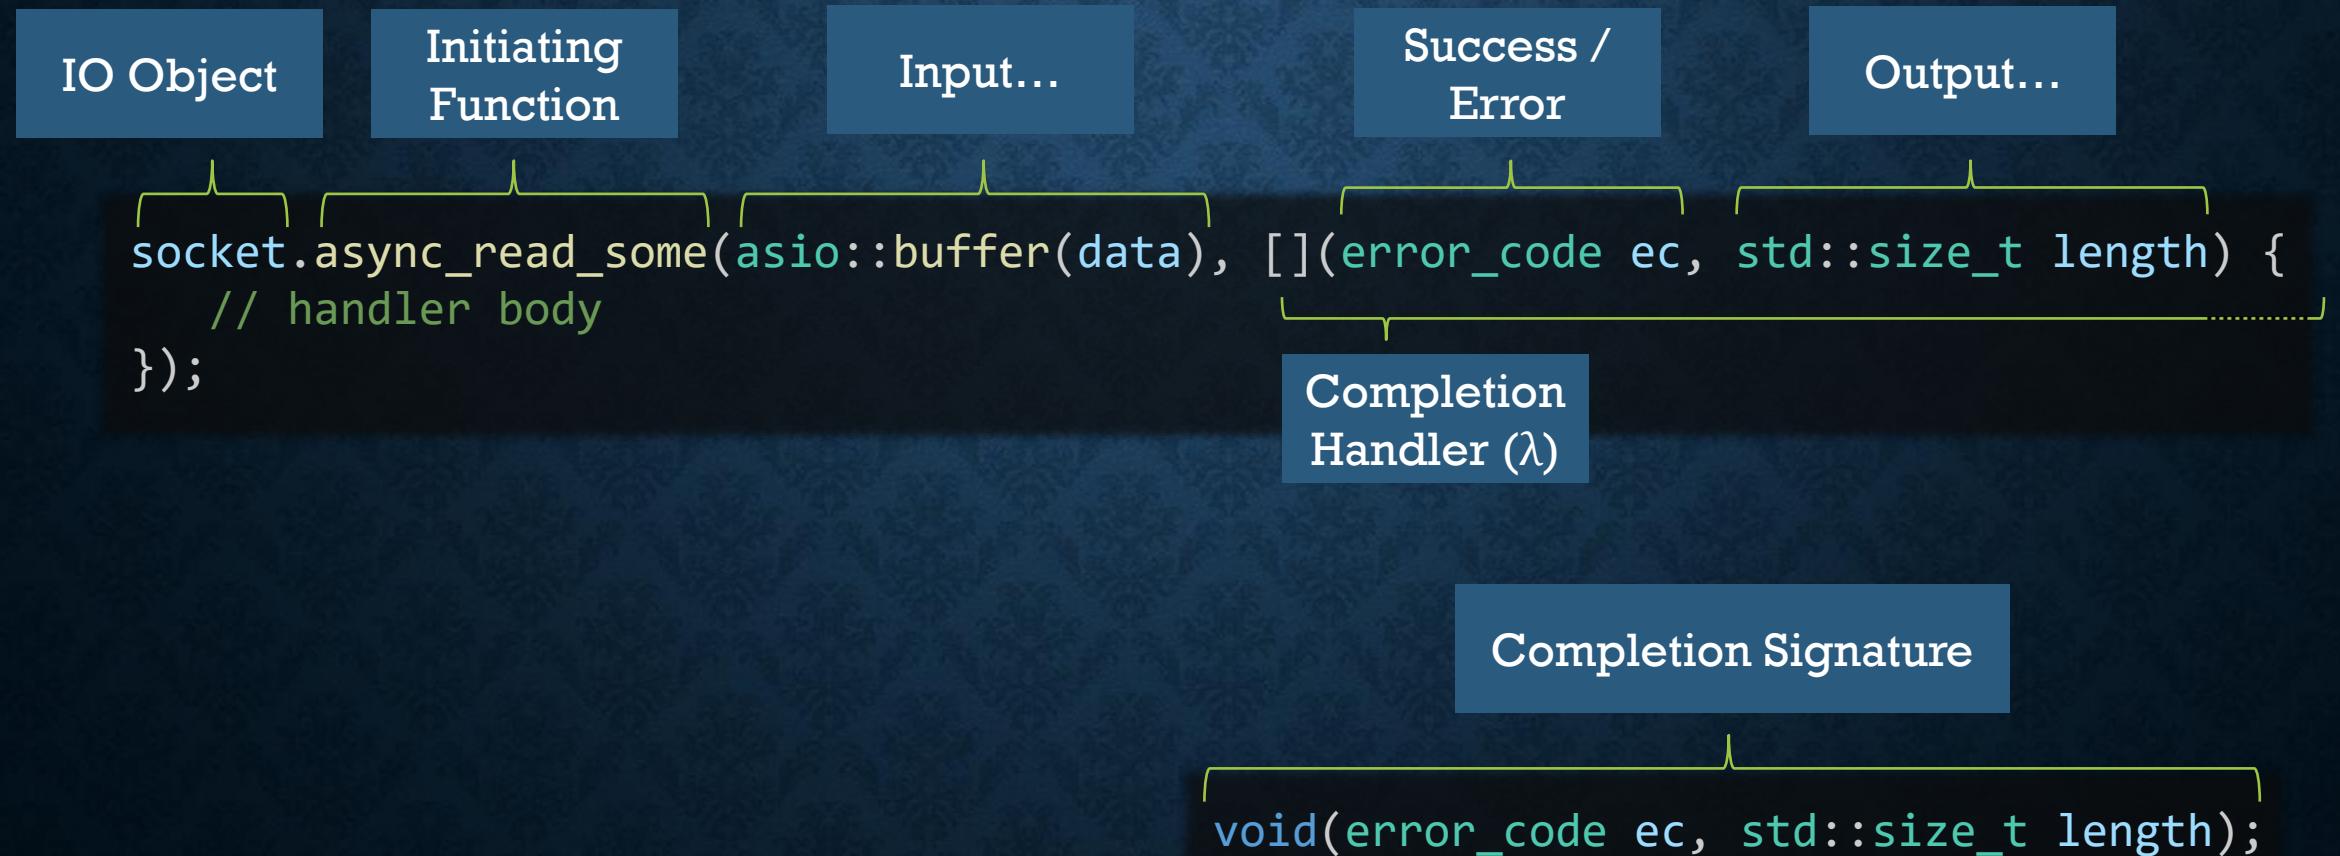
CAN WE DO BETTER?

Yes, we can.

But first let's have a closer look at how
Asynchronous Operations
are implemented in ASIO.

COMPLETION HANDLERS AND TOKENS

ASYNCHRONOUS OPERATIONS



ASYNCHRONOUS OPERATIONS

```
timer.async_wait([](error_code ec) {
    // no input, no output
});
```

```
async_write(socket_, buffer(data_, length), [](error_code ec, size_t /* length */) {
    // free function
});
```

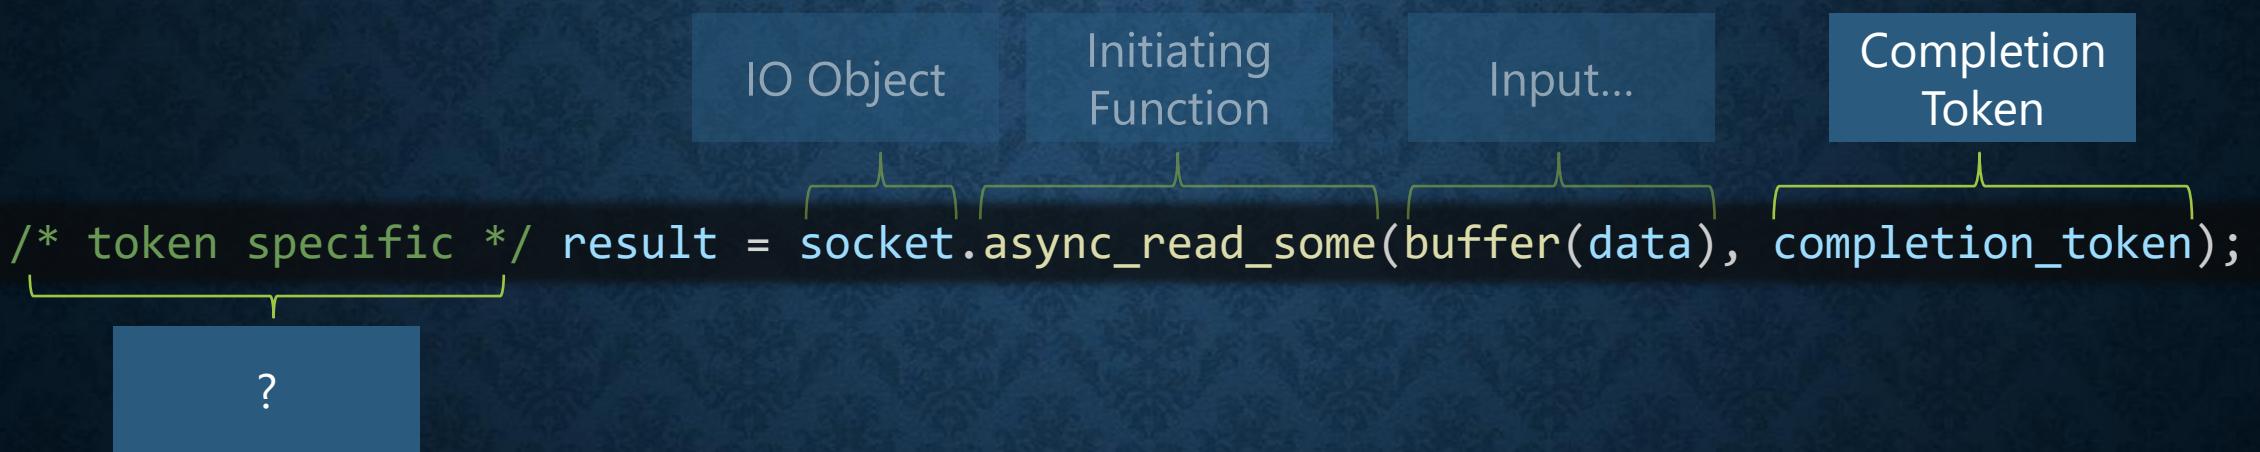
```
post([] {
    // completes unconditionally
});
```

ASYNCHRONOUS OPERATIONS

- „Initiating Function“, doesn't run anything yet but schedules for eventual completion
- “Executor” to track work and perform intermediate operations
- „Completion Handler“, invoked when completed (as per `post()`)

- But... the general async model of ASIO can do more
- A „Completion Handler“ is just one type of „Completion Token“

GENERAL INITIATING FUNCTION



COMPLETION TOKENS

Token	Description
Callable (e.g., λ)	The most common token; passes the result directly to the handler.
use_future	Returns a <code>std::future</code> representing the async result.
use_awaitable	Enables C++20 coroutines via <code>co_await</code> (returns <code>awaitable<T></code>).
deferred	Returns a deferred operation object (lazy async op), which can later be completed by calling <code>.operator()(handler)</code> or <code>operator co_await()</code> .
yield_context	Used for stackful coroutines with <code>Boost.Coroutine(spawn)</code> .
detached	Fire-and-forget: initiates the operation but ignores the result.

COMPLETION TOKEN ADAPTERS

Token Adapter	Description
<code>as_tuple</code>	Adapts the token so the completion handler receives results as a single <code>std::tuple</code> instead of throwing.
<code>redirect_error</code>	Wraps a token and redirects errors to a <code>boost::system::error_code</code> instead of throwing.
<code>bind_executor</code>	Binds an executor to a token, ensuring the handler runs on the specified executor.
<code>bind_allocator</code>	Binds an allocator to a token for handler memory allocation.
<code>cancel_{at,after}</code>	Cancel an operation if not completed before timeout (since v1.86)
<code>consign, append, prepend, ...</code>	Rarely used, prefer λ captures

TOKEN: **USE_FUTURE**

```
auto future = socket.async_read_some(buffer(data), use_future);  
size_t n = future.get(); // throws on error
```

TOKEN ADAPTER: AS_TUPLE

```
auto future = socket.async_read_some(buffer(data), as_tuple(use_future));  
auto [ec, n] = future.get();
```

?

THE deferred COMPLETION TOKEN

```
void do_read()
{
    auto self(shared_from_this());
    socket_.async_read_some(asio::buffer(data_),
                           [this, self](error_code ec, size_t length)
                           {
                               if (!ec)
                                   do_write(length);
                           });
}

void do_write(size_t length)
{
    auto self(shared_from_this());
    async_write(socket_, asio::buffer(data_, length),
                [this, self](error_code ec, size_t /* length */)
                {
                    if (!ec)
                        do_read();
                });
}
```

```
awaitable<void> session(tcp::socket socket)
{
    std::array<char, 64*1024> data;

    for (;;)
    {
        auto [ec, n] = co_await socket.async_read_some(buffer(data),
                                                       as_tuple(deferred));
        if (ec)
            break;

        std::tie(ec, n) = co_await async_write(socket, buffer(data, n),
                                                as_tuple(deferred));
        if (ec)
            break;
    }
}
```

Let's compare...

Coroutines are like Callbacks, in a specific pattern:

- The callback is transformed into the code until the next `co_await`
- Member variables become part of the coroutine frame

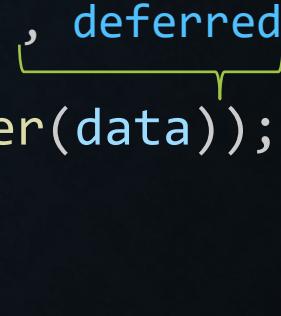
```
tcp::socket socket_;
std::array<uint8_t, 64*1024> data_;

void do_read()
{
    auto self(shared_from_this());
    socket_.async_read_some(buffer(data_),
        [this, self](error_code ec, size_t length)
    {
        if (!ec)
            do_write(length);
    });
}

void do_write(size_t length)
{
    auto self(shared_from_this());
    boost::asio::async_write(socket_, buffer(data_, length),
        [this, self](error_code ec, size_t /* length */)
    {
        if (!ec)
            do_read();
    });
}
```

```
awaitable<void> session(tcp::socket socket)
{
    std::array<char, 64*1024> data;
    for (;;)
    {
        auto [ec, n] = co_await socket.async_read_some(buffer(data),
                                                       as_tuple(deferred));
        if (ec)
            break;
        std::tie(ec, n) = co_await async_write(socket, buffer(data, n),
                                                as_tuple(deferred));
        if (ec)
            break;
    }
}
```

```
awaitable<void> session(tcp::socket socket)
{
    std::array<char, 64*1024> data;
    for (;;)
    {
        size_t n = co_await socket.async_read_some(buffer(data));
        co_await async_write(socket, buffer(data, n));
    }
}
```



default since v1.86

```
using Socket = deferred_t::as_default_on_t<tcp::socket>;  
  
awaitable<void> session(Socket socket)  
{  
    std::array<char, 64*1024> data;  
    for (;;) {  
        size_t n = co_await socket.async_read_some(buffer(data));  
        co_await async_write(socket, buffer(data, n));  
    }  
}
```

THE deferred COMPLETION TOKEN

- Makes the initiating function return an object that can be `co_await`'ed
- `deferred` is the default completion token as of Boost::Asio Version 1.86
- If the operation completes with an error, that is thrown on `co_await`
- Behavior can be modified using token adapters
- Also supports some other forms of continuations
- Here, the most interesting property is that it is `co_await`'able
- It's lazy – that means that the async operation is not started until awaited

AS-SYNC ASYNCHRONOUS OPERATIONS PROS & PROS

- Fast – a scheduler designed specifically for events is 1-2 magnitudes faster
- Scales to thousands of connections
- No synchronization required – single thread only
- Cancellation works mostly as expected
- Natural control flow (`for`, `if`, `while`, ...)
- Can use child `functions` coroutines
- Temporary storage is managed naturally (RAII)
- Lifetime of arguments is clear, ~~including the ability to safely pass temporaries~~
- Allow for Structured Programming

?

CO_SPAWN

```
class server
{
public:
    explicit server(tcp::acceptor acceptor)
        : acceptor_(std::move(acceptor))
        do_accept();
    }

private:
    void do_accept()
    {
        acceptor_.async_accept([this](error_code ec, tcp::socket socket)
        {
            if (!ec)
                std::make_shared<session>(std::move(socket))->start();
            do_accept();
        });
    }

    tcp::acceptor acceptor_;
};
```

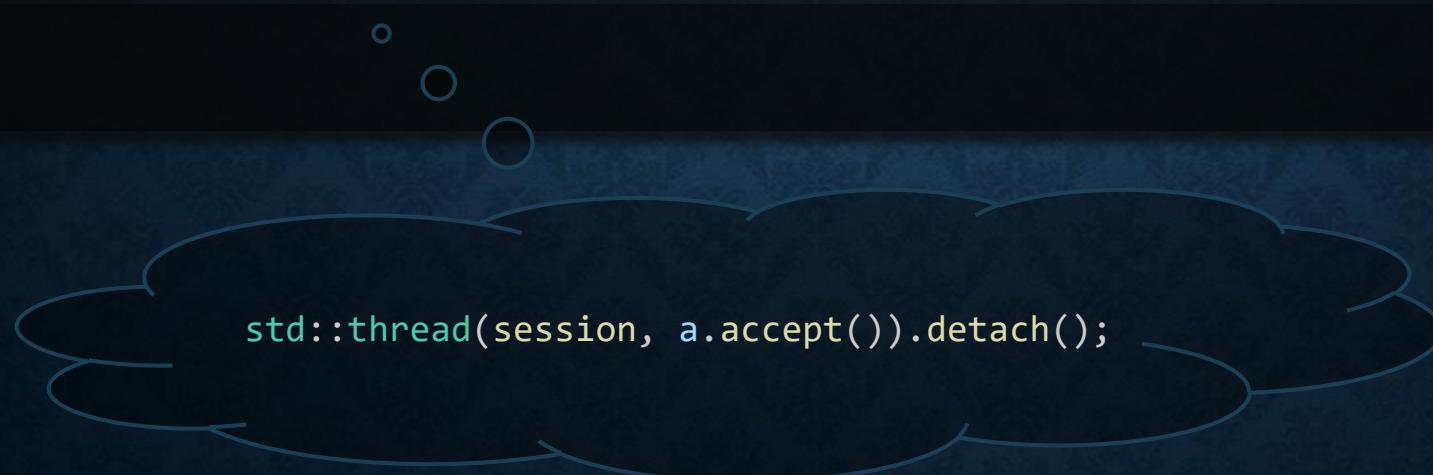
```
awaitable<void> server(tcp::acceptor acceptor)
{
    for (;;)
    {
        auto socket = co_await acceptor.async_accept();
        co_await session(std::move(socket));
    }
}
```

, deferred

What's wrong
here?

std::thread(session, a.accept()).detach();

```
awaitable<void> server(tcp::acceptor acceptor)
{
    for (;;)
    {
        auto socket = co_await acceptor.async_accept();
        co_spawn(acceptor.get_executor(), session(std::move(socket)), detached);
    }
}
```

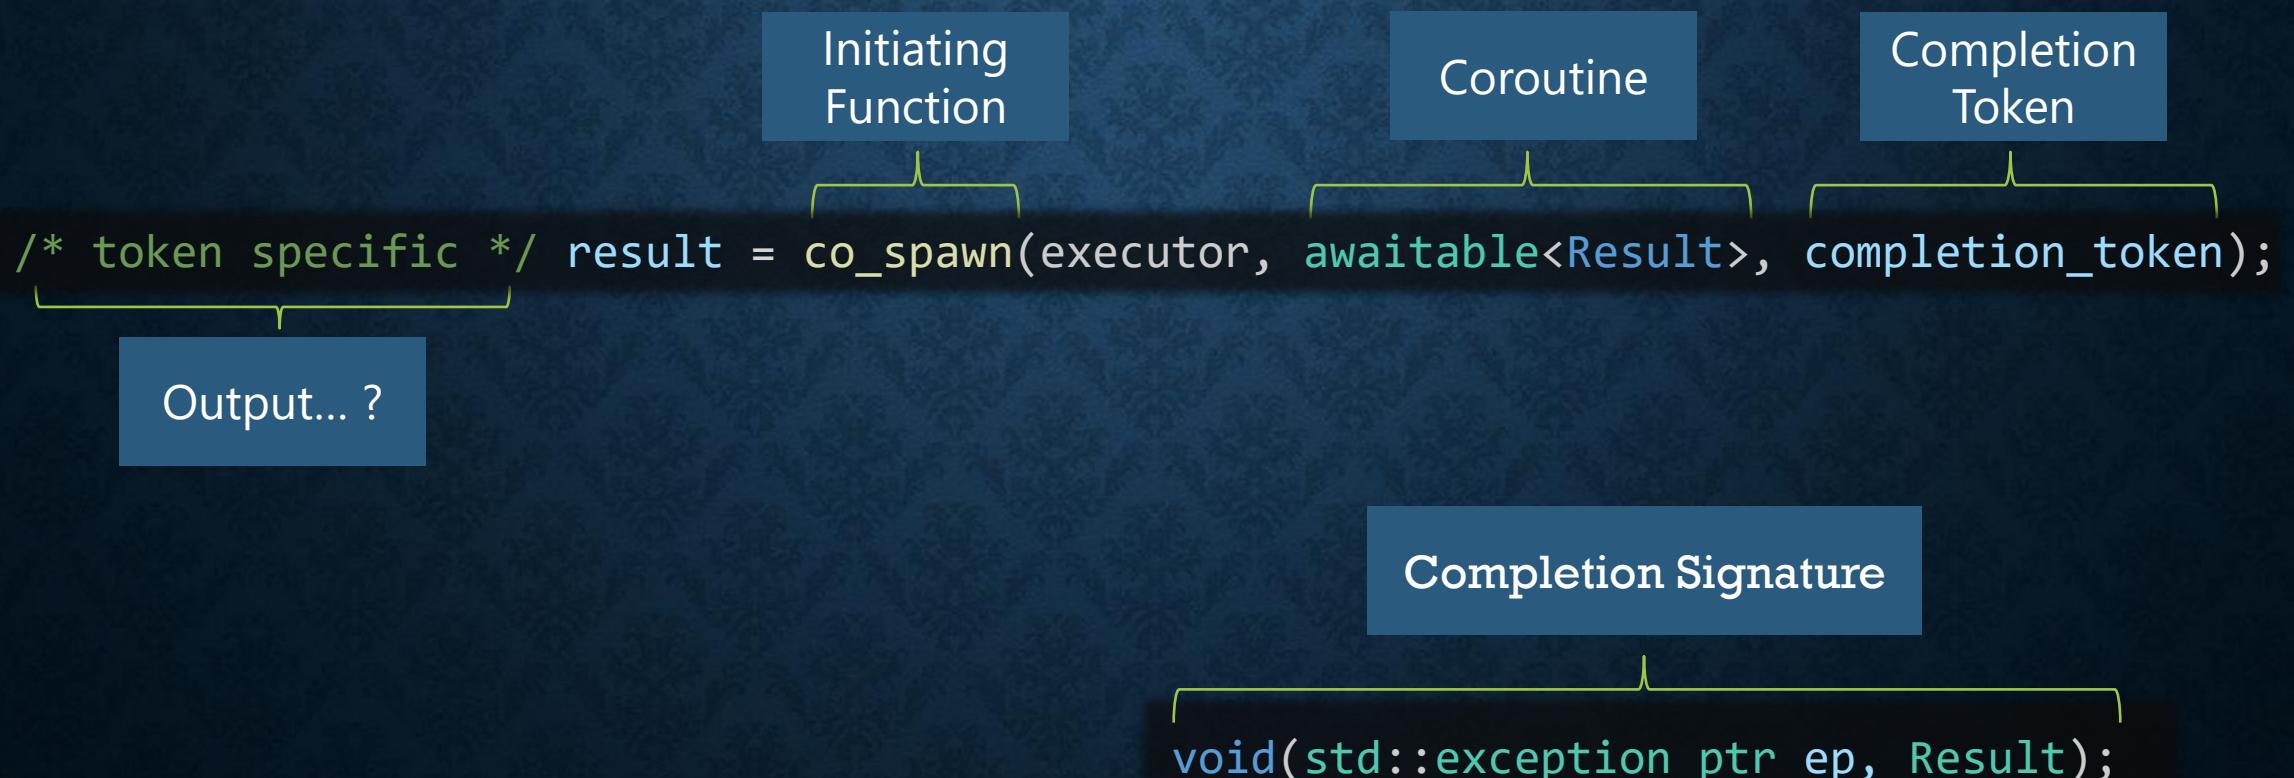


CO_SPAWN

- ASIO function, not a C++ keyword
- Creates a new coroutine-based “thread of execution”
- `co_spawn` starts an asynchronous operation
- Takes a completion token just as any other initiation function
- Unlike the other `async` operations, the completion signature is
 - `(std::exception_ptr ep, ...)` instead of
 - `(error_code ec, ...)`.

```
co_spawn(executor, session(std::move(socket)), detached);
```

CO_SPAWN



```
class server
{
public:
    server(io_context& context, tcp::endpoint endpoint)
        : acceptor_(context, endpoint) {
        do_accept();
    }

private:
    void do_accept()
    {
        acceptor_.async_accept([this](error_code ec, tcp::socket socket)
        {
            if (!ec)
                std::make_shared<session>(std::move(socket))->start();
            do_accept();
        });
    }

    tcp::acceptor acceptor_;
};
```

```
awaitable<void> server(tcp::acceptor a)
{
    for (;;)
        co_spawn(a.get_executor(),
            session(co_await a.async_accept()),
            detached);
}
```

Spawning a coroutine with the `detached` completion token is similar to what we did with the `shared_ptr` and `start()` in the `async` approach.

```
int main()
{
    io_context context;
    co_spawn(context, server({context, {tcp::v6(), 55555}}), detached);
    context.run();
}
```

```
#include <boost/asio.hpp>

using namespace boost::asio;
using ip::tcp;

awaitable<void> session(tcp::socket socket)
{
    std::array<char, 64*1024> data;
    for (;;)
    {
        size_t n = co_await socket.async_read_some(buffer(data));
        co_await async_write(socket, buffer(data, n));
    }
}

awaitable<void> server(tcp::acceptor acceptor)
{
    for (;;)
        co_spawn(acceptor.get_executor(), session(co_await acceptor.async_accept()), detached);
}

int main()
{
    io_context context;
    co_spawn(context, server({context, {tcp::v6(), 55555}}), detached);
    context.run();
}
```

from sync...

```
void session(tcp::socket socket)
{
    std::array<char, 64*1024> data;
    for (;;)
    {
        boost::system::error_code ec;
        size_t n = socket.read_some(buffer(data), ec);
        if (ec == error::eof)
            return;
        write(socket, buffer(data, n));
    }
}

void server(tcp::acceptor a)
{
    for (;;)
        std::thread(session, a.accept()).detach();
}

int main()
{
    io_context context;
    server({context, {tcp::v6(), 55555}});
}
```

... to as-sync

```
awaitable<void> session(tcp::socket socket)
{
    std::array<char, 64*1024> data;
    for (;;)
    {
        size_t n = co_await socket.async_read_some(buffer(data));
        co_await async_write(socket, buffer(data, n));
    }
}

awaitable<void> server(tcp::acceptor a)
{
    for (;;)
        co_spawn(a.get_executor(),
                 session(co_await a.async_accept()), detached);
}

int main()
{
    io_context context;
    co_spawn(context, server({context, {tcp::v6(), 55555}}),
              detached);
    context.run();
}
```

?

CONCLUSION

- ASIO does event-driven IO, thus „Inversion of Control“
- With callbacks, we lose the ability to reason about our program easily
- By introducing Coroutines, we regain the advantages of Structured Programming

ASIO & CORO

Inversion of Inversion of Control

Peter Eisenlohr
peter@eisenlohr.org
github.com/pgit/asio-coro



RESOURCES

- Boost ASIO Documentation
 - [C++20 Examples](#)
 - [P2444r0 - The Asio asynchronous model](#) (PDF)
 - https://www.boost.org/doc/libs/1_89_0/doc/html/boost_asio/reference/asynchronous_operations.html
- Blogs
 - [Notes on structured concurrency, or: Go statement considered harmful](#)
 - <https://cppalliance.org/ruben/2025/04/10/Ruben2025Q1Update.html>
 - <https://cppalliance.org/asio/2023/01/02/Asio201Timeouts.html>
 - <https://isocpp.github.io/CppCoreGuidelines/CppCoreGuidelines#cpcoro-coroutines>
- Repositories
 - <https://github.com/pgit/asio-coro> (this talk)
 - <https://github.com/anarthal/servertech-chat>

github.com/pgit/asio-coro/doc/resources.md

RESOURCES

- Using Asio with C++20 Coroutines
 - [Asynchrony with ASIO and coroutines - Andrzej Krzemieński - code::dive 2022](#)
 - [Cancellations in Asio: a tale of coroutines and timeouts - Rubén Pérez Hidalgo \(2025-06\)](#)
 - [C++20 Coroutines, with Boost ASIO in production: Frightening but awesome](#)
- Chris Kohlhoff
 - [Talking Async Ep1: Why C++20 is the Awesomest Language for Network Programming](#)
 - [Talking Async Ep2: Cancellation in depth](#)
- Structured Concurrency
 - [Writing Safer Concurrent Code with Coroutines... - Lewis Baker - CppCon 2019](#)
- Stack Overflow
 - Whatever user “sehe” says.

github.com/pgit/asio-coro/doc/resources.md